01 — Architecture

Decision: Path A (plain-Python orchestration)

The engine is a Python package `amplify/` run as a CLI:

```
python -m amplify run-weekly --week 2026-W43
```

- Each marketing agent = one Python function: system prompt (from `agents/<name>.md`) + structured input files + Mistral API call with strict `json_schema` output + schema validation.
- Orchestration is plain, readable Python. State lives in git-committed files, so a crash mid-cycle loses nothing.
- Execution environments:
  - **Manual:** UserLAnd (Kali/Ubuntu) or Termux on the founder's phone.
  - **Scheduled:** GitHub Actions cron (`.github/workflows/weekly.yml`, Thursday 06:00 EAT). The phone is never required to be awake.
- **Path B (Mistral Workflows platform) is explicitly deferred.** Migrate only when ALL of these hold: 8 weeks of pilot data, cycles failing for durability (not prompt) reasons, Path A cost exceeds quoted platform cost, platform pricing public. Migration = wrap agent functions as activities; prompts/schemas/content unchanged. Do not build Path B pieces early.

The Claims Registry — the truth layer

Single source of truth: `claims/registry.json`. Every factual statement in any content item MUST cite registry entries that exist and are unexpired.

Entry schema:

```
{
  "claim_id": "CLM-007",
  "statement": "The Dexter bot has a recorded suite of 132 tests",
  "evidence_type": "repo_commit",
  "evidence_ref": "github.com/Sami-rixx/dexter @ 59a5a6d",
  "verified_by": "samuel",
  "verified_date": "2026-10-21",
  "expires": "2027-04-21"
}
```

Evidence types: `repo_commit`, `ci_run`, `deployed_url`, `artifact`, `screenshot`, `personal_statement`.

`personal_statement` covers facts only Samuel can verify (e.g. "I study Korean", "I built this on a phone"). Samuel approves each once at registration; the entry then becomes reusable until expiry.

Enforcement: content CI rejects any item whose `claims:` list references a missing/expired entry. Unverifiable ideas are simply not said.

Signature formats (the IP)

ID	Format	Cadence	Primary platform	Purpose
FMT-01	**Phone to Product** —a real build, idea→working product, documented on the founder' s itel RS4: brief, agent setup, failures, build log, artifact	Weekly	YouTube 8–12 min	Flagship proof; founder story + product demo + AI-build evidence in one asset
FMT-02	**The Verification Files** —one unglamorous engineering truth per episode, each tied to a real repo example ("an import is not a workflow run")	Biweekly, short-form	TikTok / Shorts	Anti-hype positioning; cheapest true asset we own
FMT-03	**Blueprint in the Wild** —practical CBC/timetable/ school-ops pain points for Kenyan school owners; product second	2–3×/month	Facebook + IG	Revenue intent

Supporting (unserialized) pillars: founder journal (Korean study / scholarship journey, dated and real), build-in-public logs with receipts, Kenyan tech/AI commentary only when there is a genuine angle.

Atomization rule: ONE hero asset per cycle, decomposed into platform-native derivatives. Never four separate content efforts.

Data flow (one weekly cycle)

```
claims/ + content/calendar + agents/ + strategy/scores.json
    → Researcher → Strategist → Producer (hero + derivatives)
    → QA/Editor (rubric + claims/brand/voice lint)
    → content/review/<week>-digest.md  →  Samuel: make review
    → decisions committed → Producer revises rejected items
    → Samuel publishes manually → stats pasted → Growth Analyst
    → strategy/scores.json + Edit Loop → next week
```

Any step failure: stop, preserve logs, mark digest DEGRADED, print GitHub Issue text. Never silently skip a gate, never endless-retry.

Stack (fixed unless Samuel approves otherwise)

Layer	Choice
Language	Python 3.12–3.15
AI	Mistral API, strict JSON-schema outputs
Orchestration	Plain Python (`amplify/`)
Scheduler	GitHub Actions cron
Source control	Git + GitHub (branches, PRs, CI)
Storage	Versioned files; `analytics/marketing.sqlite` (gitignored)
Secrets	env `MISTRAL_API_KEY`; GitHub Secret for CI; never committed
Notifications	Markdown review digests; no messaging APIs in v1

Repo map

```
mobius-muse-marketing/
├── AGENTS.md                  ← build-agent constitution (root)
├── README.md
├── docs/                      ← 00..06 (this file set)
├── brand/                     ← foundation, voice-fingerprint, library, visual, banned-and-required
├── claims/registry.json + evidence/
├── strategy/                  ← audience, positioning, formats, scores.json
├── agents/                    ← runtime system-prompt files (six contracts)
├── amplify/                   ← the engine (client, agents, cycle, lint, ingest, config.yaml)
├── workflows/                 ← reserved for future Path B definitions
├── content/                   ← calendar, drafts, review, approved, published
├── campaigns/  assets/
├── analytics/                 ← marketing.sqlite (gitignored), usage.jsonl, raw/
├── schemas/                   ← content-item, rubric, per-agent output schemas
├── tests/
└── .github/workflows/         ← content-ci.yml, weekly.yml, tests.yml
```