03 — Content Item Schema & Quality Gate

The content record (one YAML file per item)

Stored at `content/drafts/<week>-<format>-<slug>.yaml`:

```yaml
content_id: 2026-W43-FMT01-EP03
format: FMT-01
campaign: w43-proof-of-work
platforms: [youtube, tiktok, instagram, facebook]
audience: [global-builders, kenya-educators]
objective: brand_trust
hook: "I built this app on a phone. Here's the receipt."
script: |
  ...
caption: ...
cta: ...
visual_brief: ...
claims: [CLM-041, CLM-042, CLM-007]
sources: ...
scores: {truth: 0, brand: 0, voice: 0, hook: 0}   # filled by QA
status: in_review    # draft | in_review | approved | rejected | published
rejection_reason: null
published_urls: {}
performance: {}
```

Rules:

- `claims` entries must exist and be unexpired in `claims/registry.json`.
- `status` transitions only: draft → in_review → (approved | rejected) → (published after manual posting). Agents may not set `approved`; only Samuel's `make review` decisions do.
- `rejection_reason` uses the taxonomy below; it feeds the learning loop.

Rejection taxonomy (reason codes)

Code	Meaning
VOICE	Doesn' t sound like Samuel
UNVERIFIED_CLAIM	Claim lacks registry evidence
OFF_BRAND	Visual/tagline/tone violation
WEAK_HOOK	Fails hook strength
WRONG_AUDIENCE	Register or platform mismatch
TOO_GENERIC	Could be any AI brand; no receipts
PERSONAL	"I' ll rewrite it myself" —no system fix implied

Rubric (QA scores each 0–100)

Dimension	Weight
Hook strength (first 1.5 s / first line)	25

Table 2 – continued

Dimension	Weight
Specificity & evidence (receipts shown)	25
Clarity	20
Brand fit	15
Platform fit	15

Auto-reject below 70. Store subscores in `scores:` so trends are measurable over time.

Content CI (runs on PRs touching content/ or claims/)

1. YAML/schema validation — malformed items never reach review.
2. **Claims lint** — every `claims:` entry exists, unexpired.
3. **Brand lint** — banned phrases ("game-changer", "revolutionary", "unlock the power"), tagline misuse ("The Inevitable Form"), off-palette visual refs, unapproved product-status words (nothing is "production-ready" unless the registry says so).
4. **Voice lint** — similarity check vs `brand/voice-fingerprint.md`; below threshold → reject.
5. Rubric presence — QA subscores present and threshold for items marked in_review.

The guardrail proof (M0 acceptance test)

CI must demonstrably catch a lie. The M0 PR includes a test draft citing non-existent claim `CLM-999`; CI fails on it; the failing run is shown in the PR, then the test draft is removed. A truth layer that has never been seen catching anything is assumed broken.

North-star metric

First-pass approval rate = share of items Samuel approves with zero edits. Target 40% at pilot end (week 4), 60% by week 12. It is the single best measure of whether the team has learned the founder's voice.