05 — Authenticity Engine & Learning Loop

The compounding moat: the system becomes more Samuel every week. Three assets, one loop.

1. Voice fingerprint — `brand/voice-fingerprint.md`

A living document describing HOW Samuel writes: typical sentence length, vocabulary, humour/earnestness balance, what he refuses to say, Kenyan vs global register cues, formatting habits. The Producer writes to it; the voice lint checks against it; the Edit Loop updates it.

Guardrail: the fingerprint is never changed by an agent directly. The Edit Loop proposes a patch as a PR; only Samuel merges it. Voice changes are founder-signed, always.

2. Voice library — `brand/voice-library/`

Real samples of Samuel's writing (posts, paragraphs, messages), tagged by context (`founder-note`, `technical-lesson`, `kenya-audience`, `global-audience`, `humorous`, `earnest`). Used as few-shot examples by the Producer. Rule: only real material enters; synthetic "brand voice" samples are forbidden (TOO_GENERIC reject at lint if detected as the source style).

3. The Edit Loop — how the team learns

Capture (automatic, at `make review`):

- Every Samuel edit to an approved draft → store before/after diff.
- Every rejection → item id + reason code + Samuel's note.

Distill (Sunday night run, `amplify editloop`):

- An agent reads the week's diffs + rejections and proposes: new voice-library samples, fingerprint patch, and Producer instruction patches — as a single PR titled `voice: week W43 learnings`.
- Samuel reviews and merges (or not). Merging IS the training signal.

This history — weeks of corrections — is the moat competitors cannot clone. Protect it: `brand/` and `content/review/decisions-*.yaml` are never deleted, only appended.

4. Performance learning — `analytics/marketing.sqlite`

Tables (minimal): `items` (content_id, format, theme, hook_text, platforms, week), `performance` (content_id, platform, metric, value, recorded_at — pasted by Samuel or ingest), `usage` (API call log).

Derived: `strategy/scores.json` — rolling scores per theme, format, hook family, audience. Written only by Growth Analyst.

Weekly memo (one question): **"What should we do more of — with evidence?"** Posted as an Issue comment. Decision rule (written down so agents don't drift): a theme earns 3 more cycles if its rolling score is above median; below-median themes are cut without sentiment. Formats are never killed before 3 full cycles.

5. Guardrail metrics (recorded weekly in the memo)

Metric	Why
First-pass approval rate	Voice learning (north star)
API cost per workflow run	Finds expensive steps
Cost per usable content item	True production efficiency
Human editing time (minutes/item)	Is automation reducing Samuel' s load?
Workflow failure rate	Reliability
Reach / engagement per published item	Distribution
Qualified inquiries (Blueprint)	Revenue connection —report honestly, including zeros

No metric is ever estimated by an agent. Unknown is a valid value.