AGENTS.md — Build-Agent Constitution

Read this file fully before writing any code in this repository. It is re-read every session. It applies to every build agent (Codex CLI, Vibe, Arena AI, or human contributor).

What this repo is

The content engine for the Möbius Muse brand: a team of six "marketing agents" (Strategist, Researcher, Evidence Keeper, Producer, QA/Editor, Growth Analyst) implemented as Python functions calling the Mistral API with strict JSON-schema outputs. Context for all design decisions lives in `docs/00-vision.md` through `docs/06-limits.md`. Read the relevant docs before starting any task.

Golden rule

AI researches, creates, checks and learns. Samuel decides what represents Möbius Muse and what gets published.

Hard rules — a PR violating any of these is rejected

1. **Secrets.** Never commit `.env`, API keys, tokens, or screenshots containing them. Never print secrets to logs. The engine reads `MISTRAL_API_KEY` from the environment only. `.env` is gitignored.
2. **Schemas.** Every marketing-agent output MUST validate against the schemas in `schemas/`. Use Mistral's strict `json_schema` response_format so malformed output is impossible by construction.
3. **Claims.** Every factual statement in any content file MUST resolve to an existing, unexpired entry in `claims/registry.json` (see `docs/01-architecture.md`). No match → the item must not reach review.
4. **No publishing.** Nothing in this repo posts, schedules, replies, messages, or buys ads — now or ever. Output stops at review digests that Samuel reads. Publication is manual, by Samuel, on each platform.
5. **Path A only.** Orchestration is plain Python in `amplify/` (see `docs/01-architecture.md`). Do not introduce new platforms, services, databases, or paid subscriptions without an Issue approved by Samuel.
6. **Branches.** `main` is stable. All work on `feat/<issue>-<slug>` branches. PRs must have green CI and updated docs.
7. **Phone constraint.** This project is developed and operated from an Android phone. No persistent daemons or long-running servers. All schedules live in GitHub Actions cron; the phone is a console (pull → review → decide → push).

Definition of done

Code + unit tests + docs updated in the same PR + `make lint` + `make test` green locally + CI green. "The agent said it's done" and "it worked once locally" do NOT count. Failures must be reproducible and logged, not silently retried into oblivion.

Verify locally (UserLAnd Kali/Ubuntu or Termux)

```
make lint        # content/schema/claims validators
make test        # unit tests, no live API calls (mocked fixtures)
make dry-run     # full weekly cycle with canned outputs, ZERO API cost
```

Generated by Kimi

Never run a live cycle (`python -m amplify run-weekly`) without Samuel's explicit instruction.

Configuration & budgets

`amplify/config.yaml` holds per-agent model choices and hard budgets (monthly token cap, per-item token cap, max retries per task = 2). Read budgets from config; never hardcode numbers. Every API call's usage is appended to `analytics/usage.jsonl`. The runner refuses to start a cycle when the monthly budget is exhausted.

When unsure

Stop and ask — comment on the Issue or PR. A halted agent is fine; a guessing agent is not. Samuel reviews from a phone; keep PRs small and diffs readable.